

Medical Insurance Cost Prediction Using Machine Learning Models.

Abstract

The aim of this project is to develop a machine learning model that predicts the cost of medical insurance for individuals based on various demographic and health-related factors. Accurate prediction of medical insurance costs can help insurance providers and individuals make informed decisions regarding coverage and financial planning. The model utilizes historical data on individuals' attributes such as age, gender, BMI, smoking habits, region, and number of dependents to train and predict the future medical insurance cost.

Introduction

This project aims to develop a machine learning model that predicts insurance charges based on individual attributes such as age, BMI, smoking habits, and region. By accurately estimating insurance charges, insurance companies can make informed decisions, improve risk assessment, and provide personalized insurance quotes to customers. The project utilizes a dataset containing relevant information and follows a standard machine learning workflow, including data preprocessing, feature selection/engineering, model training, evaluation, and fine-tuning. The ultimate goal is to enhance pricing accuracy, fairness, and customer satisfaction in the insurance industry.

Problem Statement

Insurance companies often struggle with determining appropriate insurance charges for their customers. Manual calculations may not account for the complexity of various factors affecting insurance costs, leading to potential inaccuracies and unfair pricing. This project aims to develop a predictive model that leverages machine learning techniques to accurately estimate insurance charges based on individual attributes.

Model selection

Various machine learning algorithms, such as linear regression, decision trees, random forests, or gradient boosting, will be explored to identify the most suitable model for insurance charge prediction. The model will be selected based on its ability to handle the specific characteristics of the dataset and its performance metrics.

Tools

- Numpy
- Pandas
- Matplotlib
- Seaborn
- Scikit learn

Project overview

- Importing libraries
- Reading csv file
- Data cleaning
- Data visualization
- Creation of model
- conclusion

Dataset source - Kaggle

Importing libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

Reading the csv file

In [2]:

```
df=pd.read_csv('insurance1.csv')
```

In [3]:

```
df
```

Out[3]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

Finding out how many rows and columns present in the dataset

In [4]:

```
df.shape
```

Out[4]:

```
(1338, 7)
```

Showing all the columns headings in the dataset

In [5]:

```
df.columns
```

Out[5]:

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

Datatypes of corresponding columns

In [6]:

```
df.dtypes
```

Out[6]:

```
age          int64
sex          object
bmi         float64
children     int64
smoker       object
region       object
charges     float64
dtype: object
```

Showing informations about all the features

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Categorical Features:

Sex,Smoker,Region

Checking null values for each column

In [8]:

```
df.isnull().values.any()
```

Out[8]:

```
False
```

In [9]:

```
df.isnull().sum()
```

Out[9]:

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

Checking for duplicated values

In [10]:

```
df.duplicated().sum()
```

Out[10]:

```
1
```

Dropping duplicated values

In [11]:

```
df.drop_duplicates(inplace=True)
```

In [12]:

```
df.duplicated().sum()
```

Out[12]:

```
0
```

In [13]:

```
df.head()
```

Out[13]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [14]:

```
df.tail()
```

Out[14]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

Data Analysis

Statistical Measures of the dataset

In [15]:

```
df.describe()
```

Out[15]:

	age	bmi	children	charges
count	1337.000000	1337.000000	1337.000000	1337.000000
mean	39.222139	30.663452	1.095737	13279.121487
std	14.044333	6.100468	1.205571	12110.359656
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.290000	0.000000	4746.344000
50%	39.000000	30.400000	1.000000	9386.161300
75%	51.000000	34.700000	2.000000	16657.717450
max	64.000000	53.130000	5.000000	63770.428010


Correlation

In [16]:

```
correlation_matrix = df.corr()
```

In [17]:

```
plt.figure(figsize=(6, 6))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm")
plt.title('Correlation Heatmap')
plt.show()
```



Distribution of age values

In [18]:

```
sns.set()
plt.figure(figsize=(6,6))
sns.distplot(df['age'])
plt.title("Age Distribution")
plt.show()
```




From this distribution we can conclude that most of the persons are of age between 18 to 23 and we have normal distributon from age 25 to 65.

Distribution of Gender

In [19]:

```
plt.figure(figsize=(6,6))
sns.countplot(x='sex',data=df)
plt.title('Sex Distribution')
plt.show()
```



In [20]:

```
df['sex'].value_counts()
```

Out[20]:


```
male      675
female    662
Name: sex, dtype: int64
```

Distribution of gender is almost equal

BMI Distribution

In [21]:

```
plt.figure(figsize=(6,6))
sns.distplot(df['bmi'])
plt.title("BMI Distribution")
plt.show()
```




BMI has a normal distribution. A lot of people in this particular dataset is over weighted.

Children column

In [22]:

```
plt.figure(figsize=(6,6))
sns.countplot(x='children',data=df)
plt.title('children')
plt.show()
```



In [23]:

```
df['children'].value_counts()
```

Out[23]:


```
0    573
1    324
2    240
3    157
4     25
5     18
Name: children, dtype: int64
```

More number of people who does not have any children.

Smoker column

In [24]:

```
plt.figure(figsize=(6,6))
sns.countplot(x='smoker',data=df)
plt.title('smokers')
plt.show()
```



In [25]:

```
df['smoker'].value_counts()
```

Out[25]:


```
no    1063
yes     274
Name: smoker, dtype: int64
```

Most of the people are non smokers

Region column

In [26]:

```
plt.figure(figsize=(6,6))
sns.countplot(x='region',data=df)
plt.title('Region')
plt.show()
```



In [27]:

```
df['region'].value_counts()
```

Out[27]:


```
southeast    364
southwest    325
northwest    324
northeast    324
Name: region, dtype: int64
```

Data is almost similar for all regions, a little bit more for southeast.

Distribution of charges

In [28]:

```
plt.figure(figsize=(6,6))
sns.distplot(df['charges'])
plt.title("charge Distribution")
plt.show()
```



Most of the charges are distributed in the 10000 dollars mark.

Data pre-processing

Encoding the categorical features

Encoding sex column

In [29]:

```
df.replace({'sex': {'male': 0, 'female': 1}}, inplace=True)
```

Encoding smoker column

In [30]:

```
df.replace({'smoker':{'yes':0,'no':1}},inplace=True)
```

Encoding Region column

In [31]:

```
df.replace({'region':{'southeast':0,'southwest':1,'northeast':2,'northwest':3}},inplace=True)
```

In [32]:

```
df
```

Out[32]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	0	1	16884.92400
1	18	0	33.770	1	1	0	1725.55230
2	28	0	33.000	3	1	0	4449.46200
3	33	0	22.705	0	1	3	21984.47061
4	32	0	28.880	0	1	3	3866.85520
...
1333	50	0	30.970	3	1	3	10600.54830
1334	18	1	31.920	0	1	2	2205.98080
1335	18	1	36.850	0	1	0	1629.83350
1336	21	1	25.800	0	1	1	2007.94500
1337	61	1	29.070	0	0	3	29141.36030

1337 rows × 7 columns

Splitting Features and Target

In [33]:

```
x=df.drop(columns='charges',axis=1)
```

In [34]:

```
x
```

Out[34]:

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3
...
1333	50	0	30.970	3	1	3
1334	18	1	31.920	0	1	2
1335	18	1	36.850	0	1	0
1336	21	1	25.800	0	1	1
1337	61	1	29.070	0	0	3

1337 rows × 6 columns

In [35]:

```
y=df['charges']
```

In [36]:

```
y
```

Out[36]:

```
0      16884.92400
1      1725.55230
2      4449.46200
3      21984.47061
4       3866.85520
...
1333   10600.54830
1334    2205.98080
1335    1629.83350
1336    2007.94500
1337   29141.36030
```

Name: charges, Length: 1337, dtype: float64

In [37]:

```
print(x)
```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3
...
1333	50	0	30.970	3	1	3
1334	18	1	31.920	0	1	2
1335	18	1	36.850	0	1	0
1336	21	1	25.800	0	1	1
1337	61	1	29.070	0	0	3

[1337 rows x 6 columns]

In [38]:

```
print(y)
```

```
0      16884.92400
1      1725.55230
2      4449.46200
3      21984.47061
4      3866.85520
```

```
...
1333    10600.54830
1334     2205.98080
1335     1629.83350
1336     2007.94500
1337    29141.36030
```

Name: charges, Length: 1337, dtype: float64

Splitting the data into training data and testing data

In [39]:

```
from sklearn.model_selection import train_test_split
```

In [40]:

```
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=0.2,random_state=2)
```

Shape of the test and train data.

In [41]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1069, 6)
(268, 6)
(1069,)
(268,)
```

Model Training

1. Linear Regression Model

Importing Linear Regression Model

In [42]:

```
from sklearn.linear_model import LinearRegression
```

In [43]:

```
model=LinearRegression()
```


In [44]:

```
model.fit(x_train,y_train)
```

Out[44]:

```
LinearRegression()
```

In [45]:

```
y_pred=model.predict(x_test)
```

Model Evalulation

Checking the score

In [46]:

```
model.score(x_train,y_train)
```

Out[46]:

```
0.7584123253312958
```

In [47]:

```
model.score(x_test,y_test)
```

Out[47]:

```
0.7150366419551837
```

R squared value

In [48]:

```
from sklearn.metrics import r2_score
```

In [49]:

```
r2=r2_score(y_test, y_pred)*100
```

In [50]:

```
r2
```

Out[50]:

```
71.50366419551837
```

2.Random ForestRegressor

In [51]:

```
from sklearn.ensemble import RandomForestRegressor
```

In [52]:

```
regr = RandomForestRegressor()
```

In [53]:

```
regr.fit(x_train,y_train)
```

Out[53]:

```
RandomForestRegressor()
```

In [54]:

```
y_pred=regr.predict(x_test)
```

Checking the score

In [55]:

```
regr.score(x_train,y_train)
```

Out[55]:

```
0.9776074279181237
```

In [56]:

```
regr.score(x_test,y_test)
```

Out[56]:

```
0.8339292060209827
```

R squared value

In [57]:

```
r21=r2_score(y_test, y_pred)*100
```

In [58]:

```
r21
```

Out[58]:

```
83.39292060209827
```

R squared value

In [59]:

```
r22=r2_score(y_test, y_pred)*100
```

In [60]:

```
r22
```

Out[60]:

```
83.39292060209827
```

3.Gradient BoostingRegressor

In [61]:

```
from sklearn.ensemble import GradientBoostingRegressor
```

In [62]:

```
reg = GradientBoostingRegressor()
```

In [63]:

```
reg.fit(x_train,y_train)
```

Out[63]:

```
GradientBoostingRegressor()
```

In [64]:

```
y_pred =reg.predict(x_test)
```

Checking the score

In [65]:

```
reg.score(x_train,y_train)*100
```

Out[65]:

```
91.123103331118309
```

In [66]:

```
reg.score(x_test,y_test)*100
```

Out[66]:

```
85.41167175124752
```

R squared value

In [67]:

```
r23=r2_score(y_test, y_pred)*100
```

In [68]:

```
r23
```

Out[68]:

```
85.41167175124752
```

4. Extra TreeRegressor

In [69]:

```
from sklearn.ensemble import ExtraTreesRegressor
```

In [70]:

```
regrr = ExtraTreesRegressor()
```

In [71]:

```
regrr.fit(x_train,y_train)
```

Out[71]:

```
ExtraTreesRegressor()
```

Checking the score

In [72]:

```
regrr.score(x_train,y_train)*100
```

Out[72]:

```
100.0
```

In [73]:

```
regrr.score(x_test,y_test)*100
```

Out[73]:

```
81.93650831614318
```

R square value

In [74]:

```
r234=r2_score(y_test, y_pred)*100
```

In [75]:

```
r234
```

Out[75]:

```
85.41167175124752
```

Finding the best model

In [76]:


```
results = pd.DataFrame({'Model': [ 'Linear Regression', 'Random Forest Regressor',  
                                   ' Gradient Boosting Regressor', 'Extra Tree Regressor'], 'Score': [ r2,r2,  
output_df = results.sort_values(by='Score', ascending=False)  
output_df = output_df.reset_index(drop=True)  
output_df
```

Out[76]:

	Model	Score
0	Gradient Boosting Regressor	85.411672
1	Extra Tree Regressor	85.411672
2	Random Forest Regressor	83.392921
3	Linear Regression	71.503664

In [77]:

```
plt.figure(figsize=(6,6))
plt.bar(output_df['Model'], output_df['Score'], color='blue')
plt.xlabel('Model')
plt.ylabel('Score')
plt.title('Comparison of Model Scores')
plt.ylim(0, 100)
plt.xticks(rotation=90)
plt.show()
```



conclusion

- The Gradient Boosting Regressor and Extra Tree Regressor are the top-performing models among the options listed. Both models achieved a score of 85.411672, indicating a high level of accuracy in the regression task.
- The Random Forest Regressor, with a score of 83.392921, also performed well but slightly lower than the top two models. It can still be considered a viable option, especially if the performance difference is negligible or if other factors such as interpretability or model complexity are taken into account.
- On the other hand, the Linear Regression model obtained the lowest score of 71.503664, suggesting lower accuracy compared to the other models. While Linear Regression can be a simple and interpretable baseline, it may not be the best choice for this specific task, given the available alternatives.